

```

# This code is adapted from Google Gemini.
# It retrieves all our embedding layers from the trained model and plots them according to their PCA significance.
# That is, the features are reduced to the two most important PCA components and the weights from 4 layers are plotted.
# Each color represents an important feature.

from sklearn.decomposition import PCA
from tensorflow import keras
import warnings
warnings.filterwarnings('ignore')

ALL_EMBEDDING_LAYERS = [
    'customer_embed', 'itemID_embed', 'title_embed', 'author_embed',
    'publisher_embed', 'main_topic_embed', 'subtopics_embed',
    'author_X_publisher_embed', 'author_X_customer_embed',
    'main_topic_X_customer_embed', 'subtopics_X_customer_embed',
    'title_X_main_topic_embed', 'publisher_X_main_topic_embed',
    'publisher_X_subtopics_embed', 'author_X_main_topic_embed',
    'author_X_subtopics_embed'
]
TARGET_DIM = 64

all_embeddings = []
all_source_labels = []
all_layer_names = []

# Extraction and Projection
print(f"Extracting and projecting {len(ALL_EMBEDDING_LAYERS)} embedding spaces to {TARGET_DIM} dimensions...")

for layer_name in ALL_EMBEDDING_LAYERS:
    try:
        # Retrieve the layer from the model
        embedding_layer = inference_model.get_layer(layer_name)
        weights = embedding_layer.get_weights()[0]
        embeddings = weights[1:] # Skip OOV token (index 0)

        current_dim = embeddings.shape[1]
        final_embeddings = embeddings

        # Project layers with dimension != 64D to the TARGET_DIM (64D)
        if current_dim != TARGET_DIM:
            print(f" - Projecting {layer_name}: {current_dim}D -> {TARGET_DIM}D...")

            # Create a temporary sequential model for linear projection
            projection_model = Sequential([
                Input(shape=(current_dim,)),
                Dense(TARGET_DIM, use_bias=False)
            ])

            # Project using an untrained dense layer to resize the vector space
            final_embeddings = projection_model.predict(embeddings, verbose=0)

        # Store the resulting embeddings and corresponding labels
        all_embeddings.append(final_embeddings)
        all_layer_names.append(layer_name)

        # Create labels for the plot
        labels = [layer_name] * final_embeddings.shape[0]
        all_source_labels.extend(labels)

    except Exception as e:
        # Catch layers that might not be embeddings (like 'click' or 'basket' if the list included them)
        # or layers that weren't defined in the final model structure.
        print(f"Warning: Skipping layer '{layer_name}'. Error during extraction: {e}")
        continue

# Concatenation and PCA
embedding_matrix = np.concatenate(all_embeddings, axis=0)
print(f"\nTotal embedding vectors collected: {embedding_matrix.shape[0]}")
print(f"Final vector dimension: {embedding_matrix.shape[1]}")

N_COMPONENTS = 2
print(f"Applying PCA to reduce {embedding_matrix.shape[1]}D vectors to {N_COMPONENTS}D...")

pca = PCA(n_components=N_COMPONENTS)
weights_pca = pca.fit_transform(embedding_matrix)

# Visualization
explained_variance = pca.explained_variance_ratio_.sum()

plt.figure(figsize=(12, 10))
sns.scatterplot(
    x=pca.components_[0], y=pca.components_[1],
    palette='viridis', legend=False
)

```

```

x=weights_pca[:, 0],
y=weights_pca[:, 1],
hue=all_source_labels,
s=5,
alpha=0.6,
palette=sns.color_palette("tab10", len(all_layer_names))
)

plt.title(f'PCA of ALL Learned Embedding Spaces (Total Variance Explained: {explained_variance:.1%})', fontsize=14)
plt.xlabel(f'PCA Component 1 ({pca.explained_variance_ratio_[0]:.1%})', fontsize=12)
plt.ylabel(f'PCA Component 2 ({pca.explained_variance_ratio_[1]:.1%})', fontsize=12)
plt.legend(title='Source Layer', loc='best', fontsize=8)
plt.tight_layout()
plt.show()

```

Extracting and projecting 16 embedding spaces to 64 dimensions...

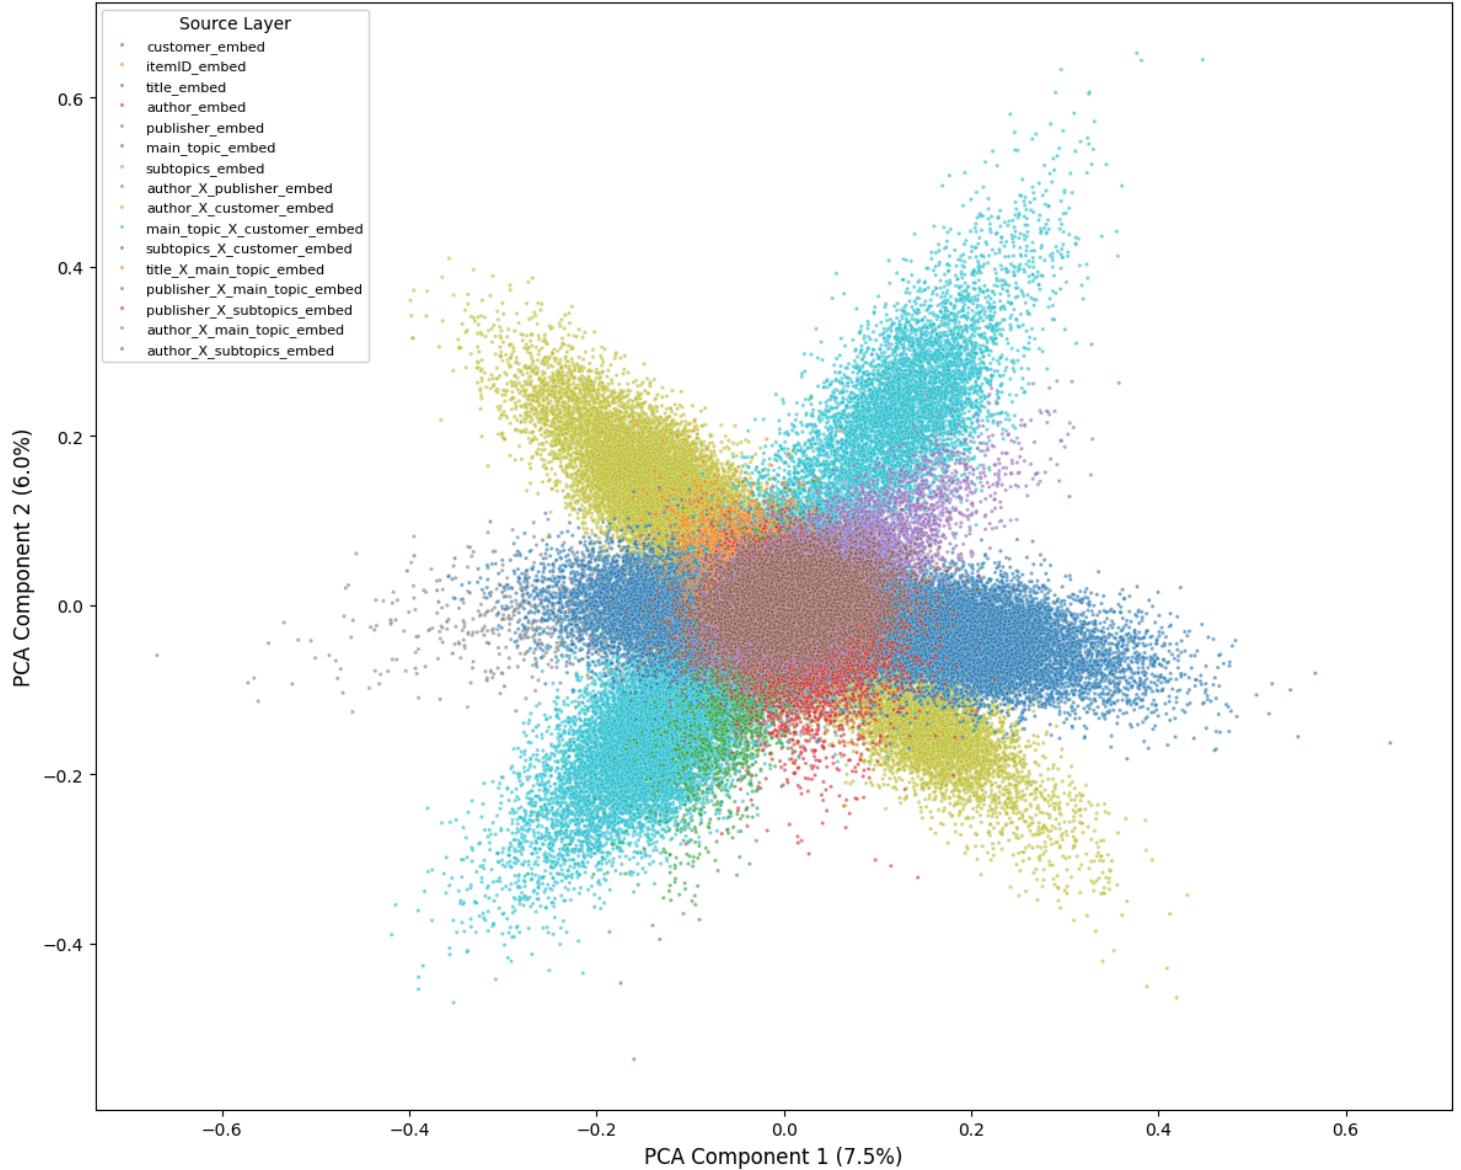
- Projecting author\_embed: 55D -> 64D...
- Projecting publisher\_embed: 25D -> 64D...
- Projecting main\_topic\_embed: 21D -> 64D...
- Projecting author\_X\_publisher\_embed: 1000D -> 64D...
- Projecting author\_X\_customer\_embed: 128D -> 64D...
- Projecting main\_topic\_X\_customer\_embed: 128D -> 64D...
- Projecting subtopics\_X\_customer\_embed: 128D -> 64D...
- Projecting title\_X\_main\_topic\_embed: 1000D -> 64D...
- Projecting publisher\_X\_main\_topic\_embed: 70D -> 64D...
- Projecting publisher\_X\_subtopics\_embed: 100D -> 64D...
- Projecting author\_X\_main\_topic\_embed: 100D -> 64D...
- Projecting author\_X\_subtopics\_embed: 100D -> 64D...

Total embedding vectors collected: 253,228

Final vector dimension: 64D

Applying PCA to reduce 64D vectors to 2D...

PCA of ALL Learned Embedding Spaces (Total Variance Explained: 13.5%)



Each cluster represents how the embedding of that feature is represented in this space, what category the embedding encodes for. For example, author\_embed (red), author\_x\_main\_topics\_embed (purple), and author\_x\_subtopics\_embed (brown) are all embedded very near / on top of each other, indicating that they're pretty similar in category. This implies that the author feature is the more important part of the crossed features (though the purple one does shoot off in a different direction, showing that the cross did teach the model some complex relationship). Dark blue and light blue are almost orthogonal, indicating that these features encode much different categories (here it's the topic vs the subtopic). Finally, the fact that they're so close together shows that these features are highly related to each other.